LeetQuiz 1 - Refined Questions

21 unique questions. Duplicates removed across all four PDFs. Based on the May 2026 exam guide.

Updated product and SQL object names have a yellow background. Substantive corrections are disclosed below each affected question. Original page references refer to archived screenshot PDFs.

Question 1 - Platform

Original: leetquiz-1.pdf, page 1

Which data lakehouse feature results in improved data quality over a traditional data lake?

A. A data lakehouse stores data in open formats.

B. A data lakehouse allows the use of SQL queries to examine data.

C. A data lakehouse provides storage solutions for structured and unstructured data.

D. A data lakehouse supports ACID-compliant transactions.

Question 2 - Transformation

Original: leetquiz-1.pdf, page 2

Which of the following describes the relationship between Gold tables and Silver tables?

A. Gold tables are more likely to contain aggregations than Silver tables.

B. Gold tables are more likely to contain valuable data than Silver tables.

C. Gold tables are more likely to contain a less refined view of data than Silver tables.

D. Gold tables are more likely to contain more data than Silver tables.

E. Gold tables are more likely to contain truthful data than Silver tables.

Question 3 - Platform

Original: leetquiz-1.pdf, page 5

Which benefit of using the Databricks Data Intelligence Platform is provided by Delta Lake?

A. The ability to manipulate the same data using a variety of languages

B. The ability to collaborate in real time on a single notebook

C. The ability to set up alerts for query failures

D. The ability to support batch and streaming workloads

E. The ability to distribute complex data operations

Question 4 - Transformation

Original: leetquiz-1.pdf, page 6

What role do notebook or other pipeline source files play in Lakeflow Spark Declarative Pipelines?

A. They only define storage locations and never contain transformations.

B. They define datasets, transformation logic, and data processing steps used by the pipeline.

C. They are used solely for visualization and reporting.

D. They are optional validation attachments unrelated to dataset definitions.

Source correction: Reworded notebook libraries as pipeline source files; modern pipelines are not restricted to notebook libraries.

Question 5 - Platform

Original: leetquiz-1.pdf, page 8

A data analysis team runs many small queries simultaneously on the same always-on SQL warehouse. Queries are slow because of concurrency. Which change addresses the queue?

A. Increase the cluster size of the SQL warehouse.

B. Increase the maximum cluster count in the SQL warehouse scaling range.

C. Turn on Auto Stop for the SQL warehouse.

D. Turn on serverless for the SQL warehouse.

Question 6 - Governance

Original: leetquiz-1.pdf, page 10

A group needs to discover objects and view table metadata in a Unity Catalog catalog named customers without reading table data. Which grant provides catalog-wide metadata discovery?

A. `GRANT SELECT ON CATALOG customers TO \`team\`;`

B. `GRANT CREATE SCHEMA ON CATALOG customers TO \`team\`;`

C. `GRANT USE SCHEMA ON SCHEMA customers.default TO \`team\`;`

D. `GRANT BROWSE ON CATALOG customers TO \`team\`;`

Source correction: Migrated the legacy database metadata-inspection question to catalog-level BROWSE. This is a privilege-model correction, not a simple rename.

Question 7 - Lakeflow Jobs

Original: leetquiz-1.pdf, page 12

When using dashboard tasks in Lakeflow Jobs, what is required for the dashboard to function properly?

A. The dashboard must be published and configured with a SQL warehouse.

B. The dashboard must use only real-time data sources.

C. The dashboard must be created using Python notebooks only.

D. The dashboard must be private to the job creator.

Question 8 - Transformation

Original: leetquiz-1.pdf, page 13

Which code removes rows where age is greater than 25 from the existing Delta table my_table and persists the change?

A. `SELECT * FROM my_table WHERE age > 25;`

B. `UPDATE my_table WHERE age > 25;`

C. `DELETE FROM my_table WHERE age > 25;`

D. `UPDATE my_table WHERE age <= 25;`

E. `DELETE FROM my_table WHERE age <= 25;`

Question 9 - Transformation

Original: leetquiz-1.pdf, page 14

A dataset in Lakeflow Spark Declarative Pipelines includes an expectation with ON VIOLATION FAIL UPDATE. What happens when a record violates the constraint?

A. Invalid records are inserted and marked invalid in the event log.

B. Invalid records are dropped and the update continues.

C. The affected flow fails and its update does not commit successfully.

D. Invalid records are inserted with an extra invalid flag.

E. Invalid records are automatically moved to a quarantine table.

Source correction: Corrected FAIL syntax and narrowed failure to the affected flow rather than claiming all independent flows stop.

Question 10 - Platform

Original: leetquiz-1.pdf, page 15

A daily dashboard refresh uses a SQL warehouse. How can the engineer minimize its active time when no queries are running?

A. Ensure each query uses a different SQL warehouse.

B. Configure the SQL warehouse as serverless without any idle policy.

C. Enable Auto Stop on the SQL warehouse.

D. Keep the SQL warehouse always running.

Source correction: Consolidated equivalent daily-refresh variants; clarified the idle-compute question without relying on legacy dashboard endpoint matching.

Question 11 - Transformation

Original: leetquiz-1.pdf, page 16

What distinguishes INSERT OVERWRITE from CREATE OR REPLACE TABLE for a Delta table?

A. INSERT OVERWRITE preserves history, whereas CREATE OR REPLACE always deletes history.

B. INSERT OVERWRITE replaces data in the existing schema; CREATE OR REPLACE TABLE can replace both data and schema.

C. Both commands always produce identical schema outcomes.

D. INSERT OVERWRITE always clears historical versions.

E. INSERT OVERWRITE replaces both data and schema, whereas CREATE OR REPLACE only replaces data.

Question 12 - Platform

Original: leetquiz-1.pdf, page 18

Engineering and analytics teams produce conflicting reports because they use siloed architectures. How can a data lakehouse help?

A. Both teams would autoscale their work as data size evolves.

B. Both teams would use the same source of truth for their work.

C. Both teams would reorganize to report to the same department.

D. Both teams would collaborate on projects in real time.

E. Both teams would respond more quickly to ad-hoc requests.

Question 13 - Governance

Original: leetquiz-1.pdf, page 19

A Unity Catalog catalog has a configured managed storage location. A new schema is created in it without MANAGED LOCATION: `CREATE SCHEMA sample_schema;` Which location is inherited for managed table storage?

A. Always dbfs:/user/hive/warehouse

B. The workspace DBFS root folder

C. The Hive metastore's spark.sql.warehouse.dir setting

D. The managed storage location configured for the parent catalog

E. A random unrelated location assigned by Databricks

Source correction: Specified a parent catalog with managed storage; schema creation itself does not create a directory containing data.

Question 14 - Governance

Original: leetquiz-1.pdf, page 20

What happens when DROP TABLE table_name removes an external Delta table?

A. The catalog registration, Delta log, history, and data files are deleted.

B. The data files are deleted but the Delta log remains.

C. The Delta log is deleted but the data files remain.

D. The catalog registration is removed; the Delta log and data files remain in external storage.

E. The data files are deleted but the catalog registration remains.

Question 15 - Transformation

Original: leetquiz-1.pdf, page 21

A DataFrame df contains product, quantity, and price. Electronics gets a 10% discount and Groceries gets 20%. Which row-level query computes the discounted total without duplicating rows? Choose one.

A. `SELECT product, quantity, price, quantity * price AS total_cost FROM df`

B. `SELECT product, quantity, price, quantity * price * CASE WHEN product = 'Electronics' THEN 0.9 WHEN product = 'Groceries' THEN 0.8 ELSE 1 END AS total_cost FROM df`

C. `SELECT product, quantity, price, quantity * price * CASE WHEN product IN ('Electronics', 'Groceries') THEN 0.8 ELSE 1 END AS total_cost FROM df`

D. `SELECT product, quantity, price, SUM(quantity * price) * CASE WHEN product = 'Electronics' THEN 0.9 WHEN product = 'Groceries' THEN 0.8 ELSE 1 END AS total_cost FROM df GROUP BY product`

E. `SELECT product, quantity, price, quantity * price * CASE WHEN product = 'Electronics' THEN 0.9 WHEN product = 'Groceries' THEN 0.8 ELSE 1 END AS total_cost FROM df UNION ALL SELECT product, quantity, price, quantity * price AS total_cost FROM df WHERE product NOT IN ('Electronics', 'Groceries')`

Source correction: Changed choose two to choose one: only B satisfies the row-level requirements; E duplicates non-discounted rows.

Question 16 - Optimization

Original: leetquiz-1.pdf, page 22

What is the primary purpose of VACUUM in Delta Lake?

A. Reorganize the physical layout of live data for faster queries.

B. Update the table schema metadata.

C. Merge small files into larger files.

D. Reclaim storage by removing unreferenced data files older than the retention threshold.

Question 17 - Platform

Original: leetquiz-1.pdf, page 25

What is a key advantage of integrating Unity Catalog with existing data systems?

A. It requires Databricks-exclusive storage systems.

B. It requires a complete overhaul of every existing catalog.

C. It operates only with non-standard SQL variants.

D. It can govern existing compatible data and storage without requiring every dataset to be physically migrated.

Source correction: Narrowed the broad original integration claim to compatible existing data and storage.

Question 18 - CI/CD

Original: leetquiz-1.pdf, page 26

What advantage does Databricks Git Folders provide over built-in notebook revision history?

A. Databricks Git Folders can revert notebook versions.

B. Databricks Git Folders is entirely contained in Databricks without remote Git.

C. Databricks Git Folders can display comments.

D. Databricks Git Folders supports multiple Git branches.

Question 19 - CI/CD

Original: leetquiz-1.pdf, page 27

Which file in a Declarative Automation Bundle allows different compute settings for dev and prod environments?

A. `config.json`

B. `databricks.yml (using targets)`

C. `.gitignore`

D. `requirements.txt`

Question 20 - Platform

Original: leetquiz-1.pdf, page 28

When INSERT OVERWRITE reloads the customer_sales Delta table, how does it affect earlier data versions?

A. It deletes all historical versions immediately.

B. It overwrites the current version and clears the transaction log.

C. It appends the data without overwriting.

D. It necessarily replaces the schema and disables time travel.

E. It overwrites current data but retains earlier versions subject to history and data retention.

Source correction: Qualified the original claim that all historical versions are preserved indefinitely.

Question 21 - Governance

Original: leetquiz-1.pdf, page 29

Which Databricks UI tool allows an authorized table owner to grant SELECT permission to other users?

A. Dashboards

B. Lakeflow Jobs

C. Databricks Git Folders

D. Catalog Explorer

Source correction: Merged Catalog Explorer and Data Explorer wording variants.