LeetQuiz 2 - Refined Questions

18 unique questions. Duplicates removed across all four PDFs. Based on the May 2026 exam guide.

Updated product and SQL object names have a yellow background. Substantive corrections are disclosed below each affected question. Original page references refer to archived screenshot PDFs.

Question 22 - Transformation

Original: leetquiz-2.pdf, page 1

Given a table events with an array column items, which query transforms each array element into a separate row?

A. `SELECT user_id, event_timestamp, event_name, ARRAY(items) AS item FROM events;`

B. `SELECT user_id, event_timestamp, event_name, items FROM events;`

C. `SELECT user_id, event_timestamp, event_name, UNNEST(items) AS item FROM events;`

D. `SELECT user_id, event_timestamp, event_name, explode(items) AS item FROM events;`

Question 23 - Platform

Original: leetquiz-2.pdf, page 2

Which stated reason for overwriting a Delta table instead of dropping and recreating it is incorrect?

A. Overwrite need not immediately delete old data files.

B. Overwrite starts a clean table history by discarding earlier audit entries.

C. Overwrite can retain earlier versions for time travel within retention limits.

D. Overwrite is atomic and does not expose a partially written table.

E. Readers can continue using a consistent snapshot while the overwrite runs.

Source correction: Clarified ambiguous clean-history wording to mean discarding history, making the incorrect claim explicit.

Question 24 - Transformation

Original: leetquiz-2.pdf, page 4

A data engineer wants a named query over two tables for use only in the current session, without storing a physical copy. Which object should they create?

A. Spark SQL table

B. Persistent view

C. Database

D. Temporary view

E. Delta table

Question 25 - Transformation

Original: leetquiz-2.pdf, page 5

Which command can combine source data into an existing Delta target with matched updates/deletes and unmatched inserts?

A. CREATE OR REPLACE TABLE

B. INSERT OVERWRITE

C. MERGE INTO

D. COPY INTO

Source correction: Removed the original guarantee that MERGE automatically handles all duplicates; source rows must have unambiguous matches.

Question 26 - Transformation

Original: leetquiz-2.pdf, page 6

A large DataFrame df has array_column containing integers. Each element must support independent filtering, sorting, and aggregation. Which approach is best?

A. Use explode to create a row for each array element.

B. Flatten all elements into a single string.

C. Always use a custom Python UDF before explode.

D. Convert arrays to strings and explode characters.

E. Treat the built-in approach and a custom UDF as equally preferable for all workloads.

Question 27 - Platform

Original: leetquiz-2.pdf, page 7

What advantage does a lakehouse provide over a traditional data lake?

A. Reduce the amount of data stored automatically.

B. Eliminate data processing.

C. Improve data quality and governance through schema enforcement.

D. Increase query complexity.

Question 28 - Governance

Original: leetquiz-2.pdf, page 15

How can a Databricks workspace segregate business units to improve data isolation and access control?

A. Create separate catalogs for business units and grant the appropriate principals access.

B. Use shared catalogs for every unit and reduce isolation.

C. Use one catalog with identical access for every unit.

D. Create overlapping catalogs and weaken access control.

Question 29 - Governance

Original: leetquiz-2.pdf, page 17

Which statement about Unity Catalog managed tables is incorrect?

A. Unity Catalog manages their metadata and data-file lifecycle.

B. Managed table data always lives outside every managed storage location configured for its schema, catalog, or metastore.

C. A CREATE TABLE statement without LOCATION normally creates a managed table.

D. DROP TABLE initiates managed-data deletion after the recovery period, rather than instantly purging all files.

E. Unity Catalog selects the managed storage path for the table.

Source correction: Reworked two conflicting incorrect choices and corrected immediate-delete wording for Unity Catalog managed tables.

Question 30 - CI/CD

Original: leetquiz-2.pdf, page 18

A Databricks Git Folders checkout was cloned from a remote Git repository. Changes have been pushed to the remote. Which operation brings them into the local checkout?

A. Clone

B. Pull

C. Merge without fetching remote changes

D. Push

Question 31 - Transformation

Original: leetquiz-2.pdf, page 19

Each product_name should have one unique category. Which query returns product names associated with more than one distinct category?

A. `SELECT product_name, COUNT(DISTINCT category) AS category_count FROM dataset GROUP BY product_name HAVING category_count > 1`

B. `SELECT product_name, MAX(category) AS max_category FROM dataset GROUP BY product_name`

C. `SELECT product_name, category FROM dataset GROUP BY product_name, category HAVING COUNT(*) > 1`

D. `SELECT product_name FROM dataset GROUP BY product_name HAVING COUNT(*) > 1`

Source correction: Specified product-name summary output; the original detail-row subquery was also a valid way to find violations and would otherwise make two options correct.

Question 32 - Transformation

Original: leetquiz-2.pdf, page 21

Which commands can query a registered Delta table from a Databricks notebook?

A.
```
%sql
SELECT * FROM table_name
```

B.
```
%python
spark.sql("SELECT * FROM table_name")
```

C.
```
%python
execute.sql("SELECT * FROM table_name")
```

D. Both A and B

E.
```
%python
delta.sql("SELECT * FROM table_name")
```

Question 33 - Transformation

Original: leetquiz-2.pdf, page 26

A pipeline needs automated validation of incoming data quality. Which tool provides declarative expectations?

A. Unity Catalog

B. Catalog Explorer

C. Delta Lake

D. Lakeflow Spark Declarative Pipelines

E. Auto Loader

Question 34 - Transformation

Original: leetquiz-2.pdf, page 27

A junior engineer uses INSERT INTO. Which command can conditionally insert only rows whose key is absent from an existing Delta target?

A. APPLY CHANGES INTO

B. MERGE INTO

C. COPY INTO

D. UPDATE

E. INSERT OR OVERWRITE

Source correction: Made key-based duplicate avoidance explicit; MERGE is not a blanket deduplication guarantee.

Question 35 - Governance

Original: leetquiz-2.pdf, page 28

How do you access customer_info in schema sales of catalog enterprise, and what does the namespace provide?

A. `SELECT * FROM enterprise.sales.customer_info;` Hierarchical organization and granular access control.

B. `SELECT * FROM sales.enterprise.customer_info;` Flat organization and broad access.

C. `SELECT * FROM customer_info;` No catalog or schema resolution is needed.

D. `SELECT * FROM enterprise_sales_customer_info;` All hierarchy is encoded in one name.

Question 36 - Platform

Original: leetquiz-2.pdf, page 29

Where is the Databricks web application hosted in the Databricks architecture?

A. Databricks Filesystem

B. Control plane

C. Customer cloud account

D. Data plane

Question 37 - Governance

Original: leetquiz-2.pdf, page 36

An engineer does not have SELECT on new_table. Assuming they can discover its metadata, where can they find its owner?

A. In the cloud object-storage file properties.

B. There is no way to find a table owner.

C. The Owner field on the table page in Catalog Explorer.

D. Any of these approaches is guaranteed.

E. Only the Permissions tab in Catalog Explorer.

Source correction: Added the metadata-discovery prerequisite; lack of SELECT alone does not guarantee or prevent owner visibility.

Question 38 - Transformation

Original: leetquiz-2.pdf, page 37

Which code updates products priced above 1000 to apply a 50% discount?

A. `SELECT price * 0.5 AS new_price FROM products WHERE price > 1000;`

B. `MERGE INTO products WHERE price > 1000 WHEN MATCHED UPDATE price = price * 0.5;`

C. `UPDATE products SET price = price * 0.5 WHERE price > 1000;`

D. `MERGE INTO products WHERE price < 1000 WHEN MATCHED UPDATE price = price * 0.5;`

E. `UPDATE products SET price = price * 0.5 WHERE price >= 1000;`

Question 39 - Platform

Original: leetquiz-2.pdf, page 38

What is a common drawback of classic job compute compared with serverless compute for jobs?

A. Job compute cannot govern data.

B. Longer startup time from provisioning compute.

C. It always has a higher total operational cost for every workload.

D. It does not support Python.

Source correction: Clarified classic job compute and avoided claiming universally higher operational cost.