LeetQuiz 4 - Refined Questions

13 unique questions. Duplicates removed across all four PDFs. Based on the May 2026 exam guide.

Updated product and SQL object names have a yellow background. Substantive corrections are disclosed below each affected question. Original page references refer to archived screenshot PDFs.

Question 64 - Transformation

Original: leetquiz-4.pdf, page 1

How do Bronze tables relate to raw data in a lakehouse?

A. They contain pre-aggregated raw data.

B. They hold a more refined view than raw data.

C. They preserve raw ingested data with a schema applied.

D. They always have less data than the raw source.

E. They always contain more accurate data than the source.

Question 65 - Ingestion

Original: leetquiz-4.pdf, page 2

Which two statements correctly describe incremental processing and its implementation?

Choose 2 answers.

A. Every run processes the entire dataset from scratch.

B. Only new or changed data since the previous run is processed, reducing work.

C. It always requires a real-time streaming technology.

D. Changes can be tracked with timestamps or CDC so only the delta is processed.

Question 66 - Transformation

Original: leetquiz-4.pdf, page 3

A pipeline reads a Bronze table, adds a calculated column, and filters invalid records into a cleaned table. Which medallion hop occurs?

A. Silver to Gold

B. Bronze to Silver

C. Raw to Bronze

D. Bronze to Gold

Source correction: Specified Bronze input and cleaned output; the original raw-table wording could otherwise be ambiguous.

Question 67 - Transformation

Original: leetquiz-4.pdf, page 4

A column json_payload contains JSON strings. Which function parses them into a Spark STRUCT using a supplied schema?

A. `to_json()`

B. `get_json_object()`

C. `from_json()`

D. `explode()`

Question 68 - Governance

Original: leetquiz-4.pdf, page 5

An authorized table owner needs to transfer table ownership to a group so its members can manage permissions. Which command is correct?

A. `TRANSFER OWNER table_name TO `group`;`

B. `GRANT OWNER ON table_name TO `group`;`

C. `ALTER TABLE table_name OWNER TO `group`;`

D. `GRANT OWNER table_name TO `group`;`

E. `ALTER OWNER ON table_name TO `group`;`

Source correction: Corrected TO syntax/principal quoting and specified an authorized owner.

Question 69 - Transformation

Original: leetquiz-4.pdf, page 13

When should CREATE OR REFRESH STREAMING TABLE be used instead of CREATE OR REFRESH MATERIALIZED VIEW in Lakeflow Spark Declarative Pipelines?

A. When the subsequent step is static.

B. When source records need incremental streaming processing.

C. Never; streaming tables are redundant.

D. For every aggregation regardless of whether a full recomputation is required.

E. When the previous step is static.

Source correction: Updated legacy live-table SQL names to modern streaming table/materialized view syntax.

Question 70 - Ingestion

Original: leetquiz-4.pdf, page 14

Auto Loader ingests JSON without schema hints or cloudFiles.inferColumnTypes. Why can numeric and boolean-looking fields be inferred as STRING?

A. The specified schema always mismatches the inferred schema.

B. For text-based formats such as JSON, Auto Loader defaults to string field inference unless type inference is enabled.

C. Auto Loader can process only string values.

D. Every field must contain a null.

E. Auto Loader cannot infer a schema.

Source correction: Clarified the default configuration; being text-based does not universally force JSON readers to produce strings.

Question 71 - Transformation

Original: leetquiz-4.pdf, page 23

A DataFrame df has URL strings in url. Which expression extracts the domain as a new column while retaining the other columns and matching an http or https URL?

A. `df.withColumn('domain', regexp_extract(col('url'), 'https?://([^/]+)', 1))`

B. `df.select(regexp_extract(col('url'), 'https?://([^/]+)', 1).alias('domain'))`

C. `df.withColumn('domain', split(col('url'), '/')[0])`

D. `df.select(split(col('url'), '/')[2].alias('domain'))`

Source correction: Original A and C both extracted the domain and B/D were valid projections. Retaining columns and regex requirement now distinguish one answer; adjusted distractor C.

Question 72 - Transformation

Original: leetquiz-4.pdf, page 24

Which workload uses a Gold table as its source?

A. Parse raw timestamps into a standard format.

B. Aggregate uncleaned raw data.

C. Remove malformed records from Bronze.

D. Query business-ready aggregated data for a dashboard.

E. Ingest a raw stream into Bronze.

Question 73 - Ingestion

Original: leetquiz-4.pdf, page 25

A table is created over CSV files. Which line fills the blank? `CREATE TABLE new_table`

`_____`

```
OPTIONS (header = "true", delimiter = "|")
LOCATION "/path/to/csv";
```

A. No provider declaration is needed.

B. USING CSV

C. FROM CSV

D. USING DELTA

E. FROM "/path/to/csv"

Source correction: Corrected the displayed path to match the stated CSV location; keep required external-location permissions in mind for Unity Catalog.

Question 74 - Ingestion

Original: leetquiz-4.pdf, page 32

How does Structured Streaming record source progress to recover after failures?

A. Checkpointing and write-ahead logs

B. Replayable sources and idempotent sinks without recorded offsets

C. Write-ahead logs and idempotent sinks without a checkpoint

D. Checkpointing and idempotent sinks, but no offset tracking

Question 75 - Ingestion

Original: leetquiz-4.pdf, page 33

A daily COPY INTO job reads files from /transactions/raw into transactions. It adds no rows because the relevant file was already loaded. Which explanation fits COPY INTO default behavior?

A. FORMAT_OPTIONS must list the file format.

B. Every load requires an explicit FILES list.

C. The previous day's file was already copied into the table.

D. PARQUET does not support COPY INTO.

E. The table must be refreshed before inserts become visible.

Source correction: Made COPY INTO explicit in the question; the original omitted the executed command.

Question 76 - Ingestion

Original: leetquiz-4.pdf, page 35

Which statement about Auto Loader's exactly-once file ingestion into Delta is correct?

A. It is implemented only through late-data watermarking.

B. File discovery metadata is tracked in RocksDB state at the checkpoint location.

C. It requires file notification mode and cannot work with directory listing.

D. It is supplied only by schema evolution.

Source correction: Clarified Delta sink and file-ingestion scope; checkpoint state alone does not guarantee every possible external sink.